

# Class 9: Neural Networks & Deep Learning

MSA 8150: Machine Learning for Analytics

---

Alireza Aghasi

Institute for Insight, Georgia State University

# Introduction

---

# Topics that Will be Covered Today

- Brief overview of linear algebra and matrix notation
- Matrix notation for linear models, especially multi-output models
- Structure of the brain
- Neural network models in matrix form
- Gradient descent technique for minimization
- NN fitting objective and (stochastic) gradient descent
- Introduction to signal processing and linear filtering
- Equipping NNs with convolutional layers
- Other variants of NNs, Recurrent NNs, Generative Adversarial Networks

# Vectors and Matrices

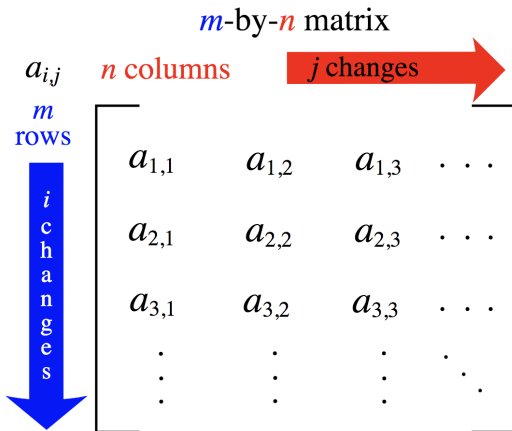
- A vector is often referred to a 1-D array of numbers

$$\mathbf{a} = \begin{bmatrix} a_1 \\ a_2 \\ \vdots \\ a_m \end{bmatrix} \in \mathbb{R}^m, \quad \text{Ex: } \mathbf{a} = \begin{bmatrix} 1.2 \\ -3 \\ 2 \end{bmatrix} \in \mathbb{R}^3.$$

- A matrix is often referred to a 2-D array of numbers

$$\mathbf{A}_{m \times n} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix}, \quad \text{Ex: } \mathbf{A} = \begin{bmatrix} 4 & -7 & 5 & 0 \\ -2 & 0 & 11 & 8 \\ 19 & 1 & -3 & 12 \end{bmatrix}.$$

# Matrix Notation



## More on Matrices

- Transpose of a matrix:

$$\mathbf{A} = \begin{bmatrix} a_{11} & a_{12} & \cdots & a_{1n} \\ a_{21} & a_{22} & \cdots & a_{2n} \\ \vdots & \vdots & \ddots & \vdots \\ a_{m1} & a_{m2} & \cdots & a_{mn} \end{bmatrix} \in \mathbb{R}^{m \times n}, \quad \mathbf{A}^T = \begin{bmatrix} a_{11} & a_{21} & \cdots & a_{m1} \\ a_{12} & a_{22} & \cdots & a_{m2} \\ \vdots & \vdots & \ddots & \vdots \\ a_{1n} & a_{2n} & \cdots & a_{mn} \end{bmatrix} \in \mathbb{R}^{n \times m}$$

- Example:

$$\mathbf{A} = \begin{bmatrix} 4 & -7 & 5 & 0 \\ -2 & 0 & 11 & 8 \\ 19 & 1 & -3 & 12 \end{bmatrix}, \quad \mathbf{A}^T = \begin{bmatrix} 4 & -2 & 19 \\ -7 & 0 & 1 \\ 5 & 11 & -3 \\ 0 & 8 & 12 \end{bmatrix}.$$

# Matrix Product

- The product of two matrices  $\mathbf{A}$ ,  $\mathbf{B}$  with compatible sizes  $n \times m$ ,  $m \times p$  is denoted by  $\mathbf{AB}$  and is of size  $n \times p$ :

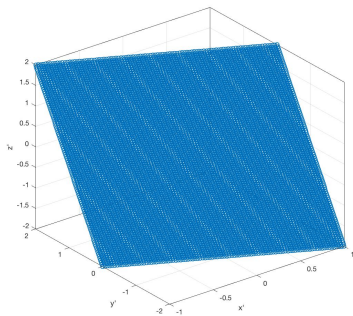
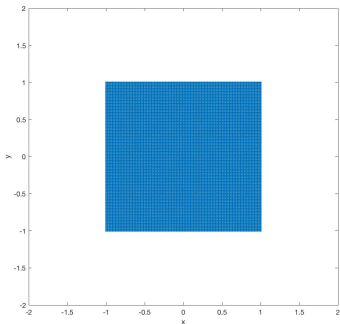
$$\begin{pmatrix} a_{11} & \cdots & a_{1m} \\ \vdots & \ddots & \vdots \\ a_{n1} & \cdots & a_{nm} \end{pmatrix} \begin{pmatrix} b_{11} & \cdots & b_{1p} \\ \vdots & \ddots & \vdots \\ b_{m1} & \cdots & b_{mp} \end{pmatrix} = \begin{pmatrix} \sum_{k=1}^m a_{1k} b_{k1} & \cdots & \sum_{k=1}^m a_{1k} b_{kp} \\ \vdots & \ddots & \vdots \\ \sum_{k=1}^m a_{nk} b_{k1} & \cdots & \sum_{k=1}^m a_{nk} b_{kp} \end{pmatrix}$$

- Example:

$$\begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & 2 \end{pmatrix} \begin{pmatrix} -2 & 0 \\ 1 & 2 \end{pmatrix} = \begin{pmatrix} -2 & 0 \\ 3 & 2 \\ 2 & 4 \end{pmatrix}$$

# Matrices as a Way of Linear Transformation

$$\begin{pmatrix} x'_1 & x'_2 & \cdots & x'_n \\ y'_1 & y'_2 & \cdots & y'_n \\ z'_1 & z'_2 & \cdots & z'_n \end{pmatrix} = \begin{pmatrix} 1 & 0 \\ -1 & 1 \\ 0 & 2 \end{pmatrix} \begin{pmatrix} x_1 & x_2 & \cdots & x_n \\ y_1 & y_2 & \cdots & y_n \end{pmatrix}$$





# Matrix Representation for the Linear Models

- For a linear model we had

$$y = b_0 + w_1x_1 + \cdots + w_px_p = w_0 + \mathbf{w}^\top \mathbf{x},$$

where

$$\mathbf{w} = \begin{pmatrix} w_1 \\ w_2 \\ \vdots \\ w_p \end{pmatrix}, \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{pmatrix}.$$

- So when we fit the model and want to evaluate it for a number of test points  $\mathbf{x}^{t_1}, \mathbf{x}^{t_2}, \dots, \mathbf{x}^{t_n}$  all we need to do is the following

$$\mathbf{y}^t = [y^{t_1}, \dots, y^{t_n}] = [b_0, b_0, \dots, b_0] + \mathbf{w}^\top \begin{pmatrix} x_1^{t_1} & x_1^{t_2} & \cdots & x_1^{t_n} \\ x_2^{t_1} & x_2^{t_2} & \cdots & x_2^{t_n} \\ \vdots & & & \vdots \\ x_p^{t_1} & x_p^{t_2} & \cdots & x_p^{t_n} \end{pmatrix}$$

# Matrix Representation for the Linear Models (Multi Response)

- For a linear model with  $m$  responses

$$y_1 = b_1 + w_{1,1}x_1 + \cdots + w_{1,p}x_p$$

$$y_2 = b_2 + w_{2,1}x_1 + \cdots + w_{2,p}x_p$$

$$\vdots$$

$$y_m = b_m + w_{m,1}x_1 + \cdots + w_{m,p}x_p$$

which can be written in the matrix form as

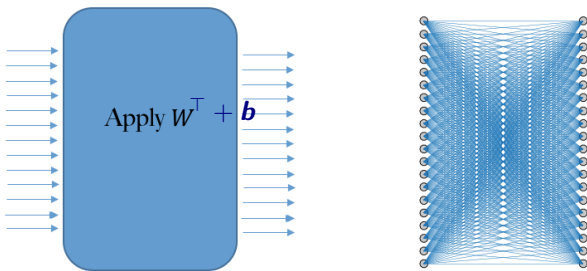
$$\mathbf{y} = \mathbf{b} + \mathbf{W}^\top \mathbf{x},$$

where

$$\mathbf{b} = \begin{pmatrix} b_1 \\ b_2 \\ \vdots \\ b_m \end{pmatrix}, \quad \mathbf{W} = \begin{pmatrix} w_{1,1} & w_{2,1} & \cdots & w_{m,1} \\ w_{1,2} & w_{2,2} & \cdots & w_{m,2} \\ \vdots & \vdots & \vdots & \vdots \\ w_{1,p} & w_{2,p} & \cdots & w_{m,p} \end{pmatrix}, \quad \mathbf{x} = \begin{pmatrix} x_1 \\ x_2 \\ \vdots \\ x_p \end{pmatrix}.$$

# Fitting Multi Response LMs

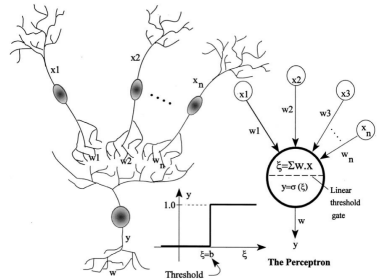
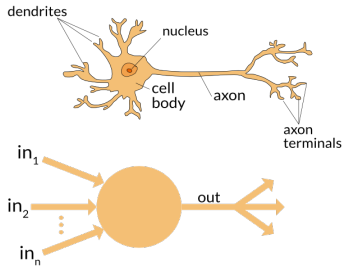
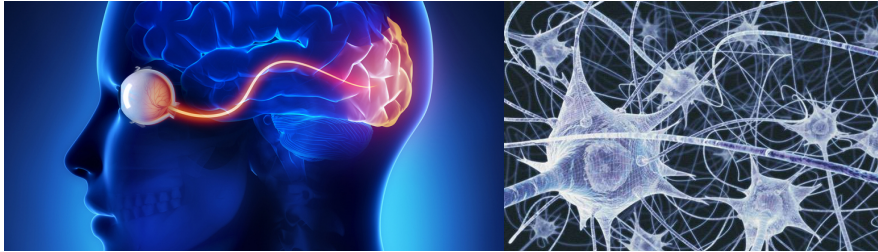
- An edge between two nodes is present when  $W_{i,j} \neq 0$



- Suppose we have the training samples  $(\mathbf{x}_1, \mathbf{y}_1), \dots, (\mathbf{x}_N, \mathbf{y}_N)$ . To fit the model to the training data, we only need to minimize the following RSS:

$$\min_{\mathbf{W}, \mathbf{b}} \sum_{i=1}^N \left\| \mathbf{y}_i - \mathbf{b} - \mathbf{W}^T \mathbf{x}_i \right\|^2$$

# Structure of the Brain



# Nonlinear Activation Applied to a Vector

- Sigmoid function:

$$\sigma(x) = \frac{e^x}{1 + e^x}$$

- When sigmoid is applied to a vector or a matrix, it applies to each component individually

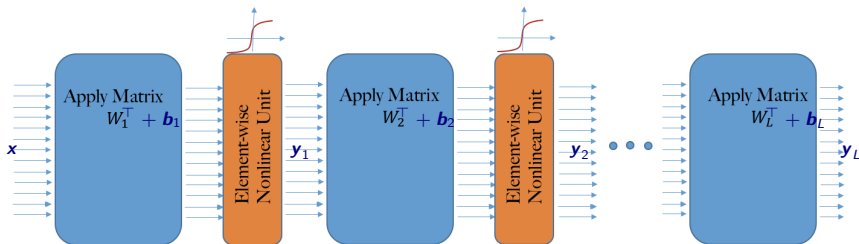
$$\sigma \left( \begin{bmatrix} 3 \\ -2 \\ -1 \end{bmatrix} \right) = \begin{bmatrix} \frac{e^3}{1+e^3} \\ \frac{e^{-2}}{1+e^{-2}} \\ \frac{e^{-1}}{1+e^{-1}} \end{bmatrix}$$

- Another widely used activation is the rectified linear unit:

$$\text{ReLU}(x) = \max(x, 0)$$

$$\text{ReLU} \left( \begin{bmatrix} 3 \\ -2 \\ -1 \end{bmatrix} \right) = \begin{bmatrix} 3 \\ 0 \\ 0 \end{bmatrix}$$

# Standard Architecture of Neural Networks



- A neural network consists of a sequence of multi-output linear units followed by nonlinear activations

$$y_1 = \sigma_1 \left( W_1^\top x + b_1 \right)$$

$$y_2 = \sigma_2 \left( W_2^\top y_1 + b_2 \right)$$

$$\vdots$$

$$y_L = \sigma_L \left( W_L^\top y_{L-1} + b_L \right)$$

# Standard Architecture of Neural Networks

- Normally all activations are taken to be identical except the last layer
- If we have regression problem, often no activation is used at the output, i.e.,

$$\mathbf{y}_L = \mathbf{W}_L^\top \mathbf{y}_{L-1} + \mathbf{b}_L$$

- For classification problems, often a soft-max unit is used at the output, i.e., for  $\mathbf{y} = [y_1, \dots, y_m]^\top$

$$\sigma_L(y_i) = \frac{e^{y_i}}{\sum_{j=1}^m e^{y_j}}, \quad i = 1, \dots, L.$$

- Example:

$$\begin{pmatrix} 0.5 \\ 1.8 \\ -2.3 \\ 0.9 \\ 0.3 \end{pmatrix} \xrightarrow{\text{soft-max}} \begin{pmatrix} 0.14 \\ 0.52 \\ 0.01 \\ 0.21 \\ 0.12 \end{pmatrix}$$

# How to Fit (Train) Neural Networks

- For the proposed architecture, we need to learn  $\mathbf{W}_1, \dots, \mathbf{W}_L$  and  $\mathbf{b}_1, \dots, \mathbf{b}_L$
- Let's first derive the function that relates  $\mathbf{x}$  to  $\mathbf{y}_L$ . Lets define

$$f_\ell(\mathbf{z}) = \sigma_\ell \left( \mathbf{W}_\ell^\top \mathbf{z} + \mathbf{b}_\ell \right),$$

then we have

$$\begin{aligned} \mathbf{y}_L &= f_L(\mathbf{y}_{L-1}) \\ &= f_L(f_{L-1}(\mathbf{y}_{L-2})) \\ &\vdots \\ &= f_L(f_{L-1}(f_{L-2}(\dots f_1(\mathbf{x}) \dots))) \end{aligned}$$

- Basically

$$\mathbf{y}_L = \mathcal{M}(\mathbf{x}), \quad \text{where} \quad \mathcal{M}(\mathbf{x}) = f_L(f_{L-1}(f_{L-2}(\dots f_1(\mathbf{x}) \dots)))$$



# How to Fit (Train) Neural Networks

- Suppose we have the training samples  $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(N)}, \mathbf{y}^{(N)})$
- For **regression** problems we normally skip an activation in the last layer and try to solve the following minimization

$$\min_{\mathbf{w}_1, \dots, \mathbf{w}_L, \mathbf{b}_1, \dots, \mathbf{b}_L} \frac{1}{N} \sum_{n=1}^N \left\| \mathbf{y}^{(n)} - \mathcal{M}(\mathbf{x}^{(n)}) \right\|^2$$

- For **classification** problems we use a soft-max in the last layer. Suppose having  $K$  classes, then  $\mathbf{y}^{(n)}$  are vectors of length  $K$ , where for each sample the corresponding index is active

$$\min_{\mathbf{w}_1, \dots, \mathbf{w}_L, \mathbf{b}_1, \dots, \mathbf{b}_L} \frac{1}{N} \sum_{n=1}^N \mathcal{H}(\mathbf{y}^{(n)}, \mathcal{M}(\mathbf{x}^{(n)}))$$

- The central objective is the cross-entropy, for  $\mathbf{y}$  and  $\mathbf{y}'$  of length  $K$ :

$$\mathcal{H}(\mathbf{y}, \mathbf{y}') = - \sum_{k=1}^K y_k \log y'_k$$

# Gradient Descent for Minimization

- We saw that our fitting problem boils down to a minimization problem

$$\min_{\mathbf{p}} \mathcal{C}(\mathbf{p})$$

in our case  $\mathbf{p}$  includes all the unknowns  $\mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L$  and  $\mathcal{C}$  is either one of the objectives in the previous slide

- Assuming  $\mathbf{p} \in \mathbb{R}^L$ , a numerical way of minimization is to start from a point  $\mathbf{p}^{(0)}$  and iteratively perform the following steps

$$\mathbf{p}^{k+1} = \mathbf{p}^{(k)} - \eta \nabla \mathcal{C} \Big|_{\mathbf{p}=\mathbf{p}^{(k)}} \quad \text{where} \quad \nabla \mathcal{C} = \begin{pmatrix} \partial \mathcal{C} / \partial p_1 \\ \partial \mathcal{C} / \partial p_2 \\ \vdots \\ \partial \mathcal{C} / \partial p_L \end{pmatrix}$$

parameter  $\eta$  is called the **learning rate**

- Let's go through a simple example to see how gradient descent works (see the MATLAB code and the next slide)

# Gradient Descent for Minimization

- Please refer to the MATLAB `gradientDescent.m` script
- Lets consider the very simple objective

$$\mathcal{C}(p_1, p_2) = (1 - p_1)^2 + (1 - p_2)^2 - 2 \exp(-3p_1^2 - 3p_2^2)$$

The gradient can be calculated as

$$\nabla \mathcal{C} = \begin{pmatrix} 2(p_1 - 1) + 12p_1 \exp(-3p_1^2 - 3p_2^2) \\ 2(p_2 - 1) + 12p_2 \exp(-3p_1^2 - 3p_2^2) \end{pmatrix}$$

- We can see that this objective has multiple local minimizers (two)
- Depending on where we start from we may land in either one
- A too small LR (learning rate) can make the minimization slow
- A too large LR can also make it slow or never converging!
- LR can affect which minimizer we converge to, but this is beyond our control

# What is Stochastic Gradient Descent?

- Recall when we had  $N$  training samples  $(\mathbf{x}^{(1)}, \mathbf{y}^{(1)}), \dots, (\mathbf{x}^{(N)}, \mathbf{y}^{(N)})$  our fitting objective was in one of the forms:

$$\min_{\mathbf{p}} \frac{1}{N} \sum_{n=1}^N \left\| \mathbf{y}^{(n)} - \mathcal{M}_{\mathbf{p}}(\mathbf{x}^{(n)}) \right\|^2 \quad \min_{\mathbf{p}} \frac{1}{N} \sum_{n=1}^N \mathcal{H}(\mathbf{y}^n, \mathcal{M}_{\mathbf{p}}(\mathbf{x}^{(n)}))$$

Here  $\mathbf{p}$  is a hyper parameter representing all our unknowns  $\mathbf{W}_1, \dots, \mathbf{W}_L, \mathbf{b}_1, \dots, \mathbf{b}_L$ .

- In other words, we are interested in objectives of the form

$$\mathcal{C}(\mathbf{p}) = \frac{1}{N} \sum_{n=1}^N \mathcal{C}_n(\mathbf{p}),$$

where  $\mathcal{C}_n$  only depends on the sample  $(\mathbf{x}^{(n)}, \mathbf{y}^{(n)})$ .

- Notice that to calculate  $\nabla \mathcal{C}$  we need to calculate  $N$  gradients

$$\nabla \mathcal{C}(\mathbf{p}) = \frac{1}{N} \sum_{n=1}^N \nabla \mathcal{C}_n(\mathbf{p})$$

# What is Stochastic Gradient Descent?

$$\nabla \mathcal{C}(\mathbf{p}) = \frac{1}{N} \sum_{n=1}^N \nabla \mathcal{C}_n(\mathbf{p})$$

- Since gradient calculation can be computationally expensive, in stochastic GD, at each minimization iteration we pick a **random** subset of all the samples  $B \subset \{1, 2, \dots, N\}$  and use it as an approximation of the gradient

$$\nabla \mathcal{C}(\mathbf{p}) \approx \frac{1}{\text{size}(B)} \sum_{n \in B} \nabla \mathcal{C}_n(\mathbf{p})$$

- If  $B$  is too small our gradient approximation may be too off!
- On the other hand large  $B$  may require a lot of gradient calculations
- Usually, after selecting the batch size,  $N_B$  we split our  $N$  data samples into  $N/N_B$  batches and in each GD iteration use one batch
- Each SGD **iteration** goes through one batch. Each **epoch** indicates going through the whole training set

# Back Propagation

- This is another terminology that you probably hear a lot in deep learning
- Recall that you had to calculate the derivative with respect to each sample and each sample function is a complicated nested function, e.g.,

$$C_n = \left\| \mathbf{y}^{(n)} - f_L(f_{L-1}(f_{L-2}(\cdots f_1(\mathbf{x}_n) \cdots))) \right\|^2, \quad f_\ell(\mathbf{z}) = \sigma_\ell \left( \mathbf{W}_\ell^\top \mathbf{z} + \mathbf{b}_\ell \right)$$

- Back propagation is simply the application of the chain rule to calculate the derivative of nested functions like  $C_n$  in terms of all the unknown parameters  $\mathbf{W}_1, \cdots, \mathbf{W}_L, \mathbf{b}_1, \cdots, \mathbf{b}_L$
- Since the actual story goes through a lot of indexing complications, let me explain things via a simple example

# Back Propagation

- You can skip this example if don't have time
- Find the derivative of the following function at  $w = 2$ :

$$f(w) = (\sin(w^2 + 1))^2$$

- **Method 1:** go through the pain of calculating the derivative and find

$$f'(w) = 4w \cos(w^2 + 1) \sin(w^2 + 1)$$

- Then plug in  $w = 2$  to get

$$f'(2) = 8 \cos(5) \sin(5)$$

# Back Propagation

- Find the derivative of the following function at  $w = 2$ :

$$f(w) = (\sin(w^2 + 1))^2$$

- **Method 2:** Notice that

$$f = g_1(g_2(g_3(w))), \text{ where } g_1(g_2) = g_2^2, g_2(g_3) = \sin(g_3), g_3(w) = w^2 + 1$$

and use the chain rule

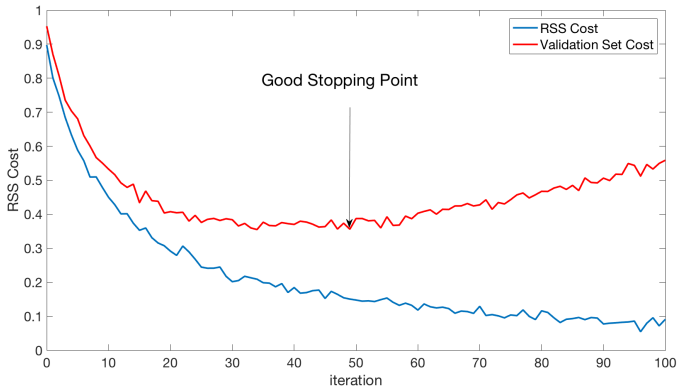
$$\frac{\partial f}{\partial w} = \frac{\partial g_1}{\partial w} = \frac{\partial g_1}{\partial g_2} \frac{\partial g_2}{\partial g_3} \frac{\partial g_3}{\partial w} = 2 \sin(5) \times \cos(5) \times 4$$

- Still want to learn more about back propagation? Try these videos:
  - <https://www.youtube.com/watch?v=Ilg3gGewQ5U>
  - <https://www.youtube.com/watch?v=tIeHLnjs5U8>



# Using a Validation Set to Control the Minimization

- As you observed in the previous slides gradient descent gradually decreases the RSS (or cross entropy cost) to find a minimizer
- One way to avoid overfitting, is to use a “**validation set**”, independent of the training set and stop the gradient descent iterations when the validation error starts to increase

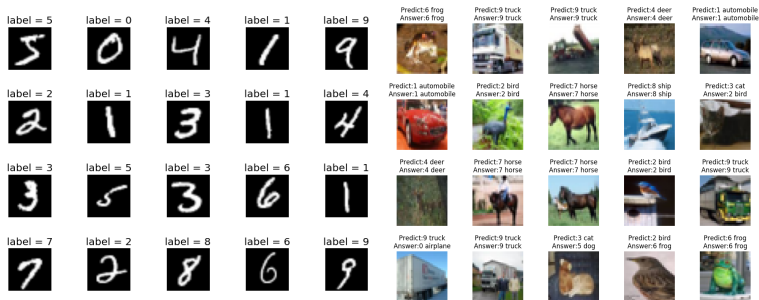


# Regularization of Neural Networks to Avoid Overfitting

- Similar to linear models there are variety of techniques to avoid overfitting in neural networks
  - $L_2$  regularizers (similar to Ridge)
  - $L_1$  regularizers (Similar to LASSO)
  - Dropout and DropConnect (which drops some of the updates in the gradient descent)
    - See video: <https://www.youtube.com/watch?v=ARq74QuavAo>
    - See papers: Paper 1, Paper 2
  - Net-Trim and compression algorithms
    - See video: <https://www.youtube.com/watch?v=WxU8dp7iYg0&t>
    - See papers and code: [Here](#)
- See one implementation of NNs **in R using the H2O package**

# Convolutional Neural Networks

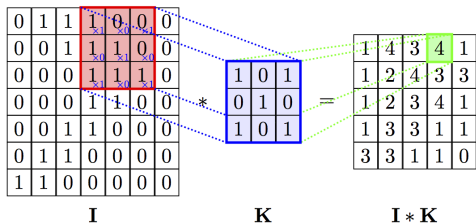
- Deep learning has shown a lot of promise in classifying images



- One immediate candidate for the features is the image pixels. But if the images are too large, this would make our problem unnecessarily large
- Convolutional neural nets are a solution to this problem

# Linear Filtering and Images

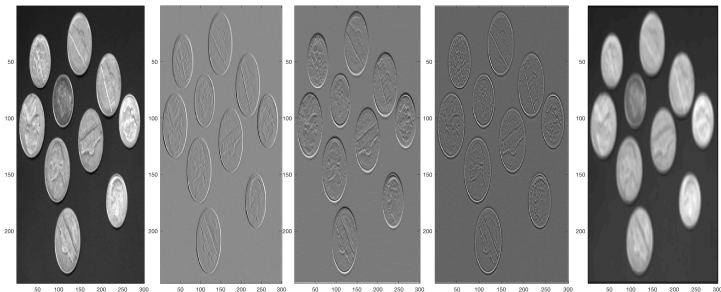
- Convolution is a linear operator widely used in image and signal processing



$$I \star K(m, n) = \sum_{i=1}^M \sum_{j=1}^N I(m-i, n-j) K(i, j)$$

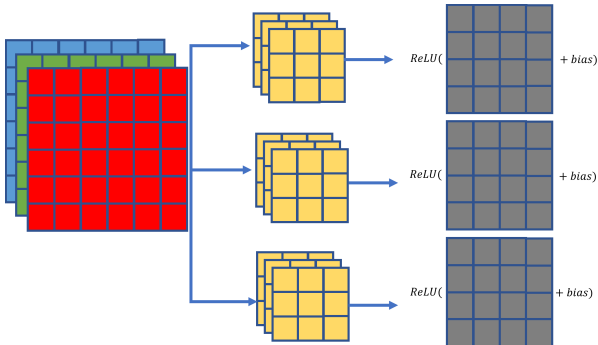
- Depending on the type of filter we pick for  $K$  the output image can have different properties (blurred, sharpened, edges detected, etc)
- See the MATLAB code `convExample.m`

## Example of Image Convolution with Different Kernels



- If the filters are selected wisely, their output can be considered as alternative features to pixels
- In a CNN, we let the neural network learn these filters! In other words, CNN wisely chooses the right features that are the best for prediction
- For color images (RGB) we can have 3D filters each filter applicable to one channel

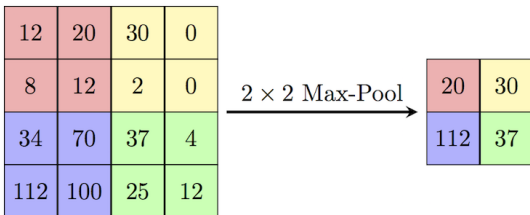
# Convolutional Layers



- We can define as many 2D or 3D convolutional filters (here 3 3D filters of size  $3 \times 3 \times 3$ )
- The total number of parameters that need to be learnt for this layer is going to be  $3 \times (27 + 1)$
- An input image of  $6 \times 6 \times 3$  is mapped to a tensor of size  $4 \times 4 \times 3$

# Max Pooling

- Is another operation that allows us to reduce the input size by taking a max operation over smaller windows across the image



- To help you better understand the topic, we will consider an implementation of the CNN that learns to read handwritten digits (MNIST data set) — see the Python code

# Recurrent Neural Networks

- While CNNs work quite promising for images, they may not be the best modeling tools for other data sets such as time series data
- For time series data and stream inputs (e.g., text analytics), **recurrent neural networks (RNNs)** are of major attention
- Remember in standard neural network the output of the hidden layer was in the form

$$\mathbf{h} = \sigma \left( \mathbf{W}_x^\top \mathbf{x} + \mathbf{b} \right)$$

- In RNNs the input is stream  $\mathbf{x}_t$  and we have another coefficient matrix that makes the current hidden output dependent on the previous one:

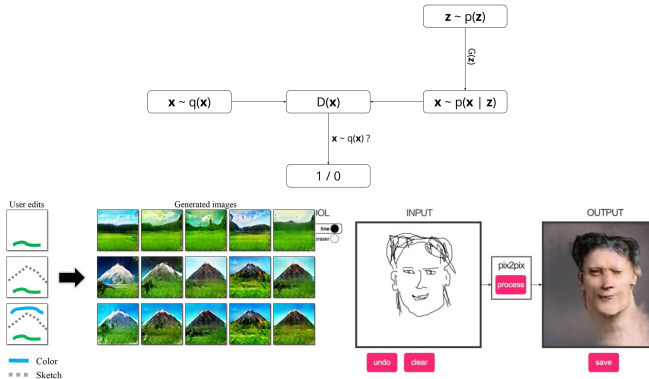
$$\mathbf{h}_t = \sigma \left( \mathbf{W}_x^\top \mathbf{x}_t + \mathbf{W}_h^\top \mathbf{h}_{t-1} + \mathbf{b} \right)$$

- To learn more and see some cool applications see:  
<https://www.youtube.com/watch?v=6niqTuYFZLQ&t=1850s>



# Generative Adversarial Networks




- Is the most recent breakthrough in machine learning started in 2015
- Basically once we pass enough samples to a GAN network, it starts to learn how to generate similar samples



- To learn more and see some interesting applications see:  
This Video ; or This Video

**Questions?**

# References

-  [http://cs231n.stanford.edu/slides/2017/cs231n\\_2017\\_lecture10.pdf](http://cs231n.stanford.edu/slides/2017/cs231n_2017_lecture10.pdf), 2016.
-  Introduction to GANs: <https://arxiv.org/pdf/1701.00160.pdf>, 2016.
-  J. Friedman, T. Hastie, and R. Tibshirani.  
**The elements of statistical learning.**  
Springer series in statistics, 2nd edition, 2009.
-  I. Goodfellow, Y. Bengio, A. Courville, and Y. Bengio.  
**Deep learning, volume 1.**  
MIT press Cambridge, 2016,  
link:<http://www.deeplearningbook.org/contents/convnets.html>.